

Table 1: V2 diagnostic: empty projected FC masks use area-weighted fallback. Complete-system results. The upper block is attributed context from OVI-MAP Table 3, not rerun here; exact external prediction protocols are not independently verified. The lower block is our paired reproduction. AP denotes the loaded released APall definition. Undefined metrics are –; reasons are recorded in cell provenance. External rows use the CVF final PDF, Table 3, checked against the arXiv v1 PDF; the supplied HTML transcription differences are retained in the source audit. The authors’ ScanNet sequence membership is not verified as CF18.

Method	Replica-8			ScanNet-CF18		
	AP↑	AP50↑	mIoU↑	AP↑	AP50↑	mIoU↑
<i>Reported in OVI-MAP (source values; not our measurements)</i>						
Mask3D + OpenMask3D	3.2	4.5	18.7	5.1	8.0	8.6
Segment3D + OpenMask3D	9.0	15.2	17.3	2.3	3.9	4.5
OVO-SLAM	9.1	17.5	24.9	5.5	12.6	14.6
<i>Paired evaluation under our fixed released-code protocol</i>						
OVI-MAP (reproduced, A0)	8.63	21.48	27.24	7.18	15.72	16.38
Ours: refinement only (A1)	11.74	24.50	29.69	8.31	17.96	19.00
Ours: refinement + recovery (A3)	12.39	26.06	30.27	8.23	17.83	19.02

Table 2: V2 diagnostic: empty projected FC masks use area-weighted fallback. Controlled ablations on the same geometry. E is fixed D2 refinement; R is independent single-view FC recovery. A0–A3 form the $E \times R$ design. A4 uses the same G1 view with native recognition. A5 uses the same incumbent FC evidence, with unknown output where unavailable; it has no native fallback. Undefined metrics are –; reasons are recorded in cell provenance. A5 retains Native-derived incumbent support; FC-only denotes its conditional readout.

Configuration	Existing	Recovery	Replica-8			ScanNet-CF18		
			AP	AP50	mIoU	AP	AP50	mIoU
A0: native	Native	None	8.63	21.48	27.24	7.18	15.72	16.38
A1: E only	D2	None	11.74	24.50	29.69	8.31	17.96	19.00
A2: R only	Native	G1-FC	9.44	23.69	27.96	7.11	15.61	16.44
A3: E + R	D2	G1-FC	12.39	26.06	30.27	8.23	17.83	19.02
A4: native recovery	D2	G1-Native	11.74	24.50	29.82	8.30	17.93	18.99
A5: FC-only	FC	G1-FC	10.94	21.88	23.72	8.32	17.39	17.94

Table 3: G1/G3 latency is remeasured on all eight scenes with the resident model and no feature/view/result cache; U2 is the unchanged control. V2 diagnostic: empty projected FC masks use area-weighted fallback. Recovery on Replica-8. n/N is recovered owners over the same prelocked candidate pool, not true positives. TP/FP are added released-matcher score entries at IoU .50. Time is feature-cache-cold incremental recovery with the model resident, averaged over all eight scenes; model loading, mapping and evaluation are separate. Undefined metrics are –; reasons are recorded in cell provenance. † marks additional mixed old/new tied score entries with unresolved attribution; displayed TP/FP counts contain only definite added entries.

Entry	Max views	n/N	TP50	FP50	AP	mIoU	Time (s)
None (A1)	0	0/53	0	0	11.74	29.69	0
Archived U2	1	9/53	3	2	12.39	30.17	6.61
Independent G1 (A3)	1	42/53	3	3	12.39	30.27	45.58
Independent G3	3	42/53	3	4	12.10	30.18	45.63