

Runtime and readout evidence: layout preview

These are templates, not measured results. The main paper keeps three tables; these two compact blocks are intended for the supplement. Accuracy is reused only with verified prediction parity. Runtime is measured in the new execution task.

Incumbent readout	Recovery	Replica-8			ScanNet-CF18		
		AP	AP50	mIoU	AP	AP50	mIoU
Native (A2)	Shared G1	—	—	—	—	—	—
FC-only (A5)	Shared G1	—	—	—	—	—	—
Grouped (A3)	Shared G1	—	—	—	—	—	—

Table S7: Readout comparison with identical recovered instances. These rows reuse same-version Table 2 outputs. They are not a backbone-only ablation.

Incremental recovery stage (wall s/scene)	U2 control	G1 reference	G1 selected
Registry and geometry preparation	—	—	—
View construction (depth I/O, ray queries, selection)	—	—	—
RGB and preprocessing	—	—	—
Recognition (encoding, pooling, classification)	—	—	—
Output and bookkeeping	—	—	—
Other / final synchronization	—	—	—
Complete incremental wall time	—	—	—

Table S8: Same-device recovery latency over eight scenes and two repetitions. The model and common map/evidence are resident; recovery features, views, BVH and results are not reused across calls. GPU service diagnostics are reported separately and must not be added to wall time. No whole-mapper FPS is inferred.